

CS 33, Fall 2020

Malloc Lab: Implementing a Dynamic Memory Allocator

Due: 12/2/2020

Mrunal Patel is the lead TA for this assignment.

1 Introduction

In this lab you will be writing a dynamic memory allocator for C programs, i.e., your own version of the malloc and free routines. You are encouraged to explore the design space creatively and implement an allocator that is correct, efficient and fast.

2 Logistics

You must work alone for this project. Any clarifications and revisions to the assignment will be posted on the course Web page.

3 Hand Out Instructions

You can download the lab files at:

<https://polyarch.github.io/cs33/labs/malloclab-fall20.pdf>

Start by copying `malloclab-handout.tar` to a directory in which you plan to do your work. Then give the command: `tar xvf malloclab-handout.tar`. This will cause a number of files to be unpacked into the directory. The only file you will be modifying and handing in is `mm.c`. The `mdriver.c` program is a driver program that allows you to evaluate the performance of your solution. Use the command `make` to generate the driver code and run it with the command `./mdriver -V`. The `-V` flag displays helpful summary information. You may ignore the warnings generated when compiling `mdriver`.

Looking at the file `mm.c` you'll notice a C structure into which you should insert the requested identifying information. **Do this right away so you don't forget.**

When you have completed the lab, you will hand in only one file (`mm.c`), which contains your solution. To hand in your lab, use the command `make handin`. You can hand in your work as many times as you would like, but only the most recent hand-in will be graded.

4 How to Work on the Lab

Your dynamic storage allocator will consist of the following four functions, which are declared in `mm.h` and defined in `mm.c`.

```

int  mm_init(void);
void *mm_malloc(size_t size);
void mm_free(void *ptr);
void *mm_realloc(void *ptr, size_t size);

```

The `mm.c` file we have given you implements an implicit list version of `malloc`. Using this as a starting place, modify these functions (and possibly define other private static functions), so that they obey the following semantics:

- `mm_init`: Before calling `mm_malloc` or `mm_free`, the application program (i.e., the trace-driven driver program that you will use to evaluate your implementation) calls `mm_init` to perform any necessary initializations, such as allocating the initial heap area. The return value should be `-1` if there was a problem in performing the initialization, `0` otherwise. Any global variables you use in your allocator must be (re-)initialized here.
- `mm_malloc`: The `mm_malloc` routine returns a pointer to an allocated block payload of at least `size` bytes. The entire allocated block should lie within the heap region and should not overlap with any other allocated block.

Since the `libc malloc` always returns payload pointers that are aligned to 8 bytes, your `malloc` implementation should do likewise and always return 8-byte aligned pointers.

- `mm_free`: The `mm_free` routine frees the block pointed to by `ptr`. It returns nothing. This routine is only guaranteed to work when the passed pointer (`ptr`) was returned by an earlier call to `mm_malloc` and has not yet been freed.
- `mm_realloc`: You do not need to modify the `mm_realloc` routine. Please leave it as it is so that the lab can compile without any issues.

These semantics match the semantics of the corresponding `libc malloc` and `free` routines. Type `man malloc` in the terminal for complete documentation.

4.1 The `block_t` struct

In `mm.c` you will find a few structs and typedefs. These are provided to you to help you manipulate the blocks in the heap easily. You can cast an address to be a pointer of type `block_t*` and then easily deference it using memorable field names. The struct is provided to you as a starting point, you are not required to use it and you are also allowed to change the struct.

Accessing hints (assuming you have a `block_t* block`):

- Accessing the allocated bit: `block->allocated;`
- Accessing the block size: `block->block_size;`
- Getting the payload address: `block->body.payload;`
- Accessing the next pointer for a free block: `block->body.next;`
- Accessing the prev pointer for a free block: `block->body.prev;`

```
typedef struct {
    uint32_t allocated : 1;    // allocated bit
    uint32_t block_size : 32; // block size
    /* unused padding (4 bytes) */
    union {
        struct {
            void* next;        // next free block pointer
            void* prev;        // prev free block pointer
        };
        char payload[0];       // pointer to the payload
    } body;
} block_t;
```

This struct is a fairly complicated struct to understand as there is a lot going on. The first two fields are bit fields. Together, they will occupy 32 bits of space. You can easily access each group of bits in your C program using this method. The LSB (i.e. the allocated bit) is at the top of the struct. The struct contains a union which represents the body of the block. The body can be a payload, if the block is allocated, or it can (for an explicit list) be a pair of pointers, if the block is free. The struct for the pointers is an anonymous struct, you cannot instantiate this struct. The union is going to be 8 byte aligned and so there are 4 padding bytes between the block size and the body.

It's worth noting how the payload field works. When you declare an array, the name is a label that represents the starting address of the array. We use this cleverly to declare an array of size 0 effectively creating an easy way to access the address of that particular location in the struct.

5 Heap Consistency Checker

Dynamic memory allocators are notoriously tricky beasts to program correctly and efficiently. They are difficult to program correctly because they involve a lot of untyped pointer manipulation. You will find it very helpful to write a heap checker that scans the heap and checks it for consistency.

Some examples of what a heap checker might check are:

- Is every block in the free list marked as free?
- Are there any contiguous free blocks that somehow escaped coalescing?
- Is every free block actually in the free list?
- Do the pointers in the free list point to valid free blocks?
- Do any allocated blocks overlap?
- Do the pointers in a heap block point to valid heap addresses?

Your heap checker will consist of the function `int mm_check(void)` in `mm.c`. It will check any invariants or consistency conditions you consider prudent. It returns a nonzero value if and only if your heap is consistent. You are not limited to the listed suggestions nor are you required to check all of them. You are encouraged to print out error messages when `mm_check` fails.

This consistency checker is for your own debugging during development. When you submit `mm.c`, make sure to remove any calls to `mm_check` as they will slow down your allocator.

6 Memory System Library

The `memlib.c` package simulates the memory system for your dynamic memory allocator. You can invoke the following functions in `memlib.c`:

- `void *mem_sbrk(int incr)`: Expands the heap by `incr` bytes, where `incr` is a positive non-zero integer and returns a generic pointer to the first byte of the newly allocated heap area. The semantics are identical to the Unix `sbrk` function, except that `mem_sbrk` accepts only a positive non-zero integer argument (i.e. you cannot decrement the heap break using `mem_sbrk`).
- `void *mem_heap_lo(void)`: Returns a generic pointer to the first byte in the heap.
- `void *mem_heap_hi(void)`: Returns a generic pointer to the last byte in the heap.
- `size_t mem_heapsize(void)`: Returns the current size of the heap in bytes.
- `size_t mem_pagesize(void)`: Returns the system's page size in bytes (4096 on Linux systems).

7 The Trace-driven Driver Program

The driver program `mdriver.c` in the `malloclab-handout.tar` distribution tests your `mm.c` package for correctness, space utilization, and throughput. The driver program is controlled by a set of *trace.rep* files that are included in the `traces` directory in the `malloclab-handout.tar` distribution. Each trace file contains a sequence of allocate and free directions that instruct the driver to call your `mm_malloc` and `mm_free` routines in some sequence. The driver and the trace files are the same ones we will use when we grade your `handin mm.c` file.

The driver `mdriver.c` accepts the following command line arguments:

- `-t <tracedir>`: Look for the default trace files in directory `tracedir` instead of the default directory defined in `config.h`.
- `-f <tracefile>`: Use one particular `tracefile` for testing instead of the default set of `tracefiles`.
- `-h`: Print a summary of the command line arguments.
- `-v`: Verbose output. Print a performance breakdown for each `tracefile` in a compact table.

- `-V`: More verbose output. Prints additional diagnostic information as each trace file is processed. Useful during debugging for determining which trace file is causing your malloc package to fail.

8 Programming Rules

- You should not change any of the interfaces (definitions of `mm_malloc`, `mm_free`, `mm_realloc`, `mm_init` in `mm.c`).
- You should not invoke any memory-management related library calls or system calls. This would exclude the use of `malloc`, `calloc`, `free`, `realloc`, `sbrk`, `brk` or any variants of these calls in your code.
- You are not allowed to define any global or static compound data structures such as arrays, structs, trees, or lists in your `mm.c` program. However, you *are* allowed to declare global scalar variables such as integers, floats, and pointers in `mm.c`. You are allowed to declare data structures globally (i.e. typedef a struct) and encouraged to cast addresses to structs you find useful.
- For consistency with the `libc` malloc package, which returns blocks aligned on 8-byte boundaries, your allocator must always return pointers that are aligned to 8-byte boundaries. The driver will enforce this requirement for you.

9 Evaluation

You will receive **zero points** if you break any of the rules or your code is buggy and crashes the driver. Otherwise, your grade will be calculated as follows:

- *Performance*. Two performance metrics will be used to evaluate your solution:
 - *Space utilization*: The peak ratio between the aggregate amount of memory used by the driver (i.e., allocated via `mm_malloc` but not yet freed via `mm_free`) and the size of the heap used by your allocator. The optimal ratio equals to 1. You should find good policies to minimize fragmentation in order to make this ratio as close as possible to the optimal.
 - *Throughput*: The average number of operations completed per second. We use Kop/s (thousands of operations per second) as the throughput.

The driver program summarizes the performance of your allocator by computing a *performance index*, P , which is based on two factors:

1. **Scaled Utilization - U_s** : This is calculated by first calculating the weighted average of the utilizations (U_{avg}) across traces. Some traces are worth more than others. The formula for scaling is:

$$U_s = \frac{U_{avg} - 0.6}{0.4}$$

2. **Weighted Harmonic Mean of Throughput - T :** The harmonic mean is used to understand averages for rates. Throughput is a rate. We use a weighted harmonic mean to emphasize some traces.

The *performance index* is calculated as follows:

$$P = \frac{U_s^2 \times T}{100}$$

This index is supposed to motivate you to try to maximize throughput with minimal impact on utilization. To receive a good score, you must achieve a balance between utilization and throughput.

- *Correctness.* Your implementation should be correct on all trace files. Use `mdriver -V` to make sure you have valid implementations.
- Documentation of your code.
 - Your code should be decomposed into functions and use as few global variables as possible.
 - Your code should begin with a header comment that describes the structure of your free and allocated blocks, the organization of the free list, and how your allocator manipulates the free list. Each function should be preceded by a header comment that describes what the function does.
 - Each subroutine should have a header comment that describes what it does and how it does it.
 - Your heap consistency checker `mm_check` should be thorough and well-documented.

10 Extra credit

Extra credit will be awarded to encourage hard work to create the fastest implementation.

- The top 20% students with the best performance score will receive a 10% bonus.
- The top 10% students with the best performance score will receive a 20% bonus.

A scoreboard is up at <http://lrxsrv06.seas.ucla.edu:15213/scoreboard>

11 Handin Instructions

Use the command `make handin` to submit your work. You may submit your solution as many times as you wish up until the due date. Only the last version you submit will be graded.

Because the performance of your code may vary slightly between runs, *we will be running your submission five (5) times and selecting the best performance out of those for your final grade.*

When testing your files locally, make sure to use one of the class machines. This will insure that the grade preview you get from `mdriver -V` is representative of the grade you will receive when you submit your solution.

Running `mdriver` will print out the unique hash for your lab. This ID identifies your submission on the scoreboard.

12 Hints

- *Use the `mdriver -f` option.* During initial development, using tiny trace files will simplify debugging and testing. There are two such trace files (`short1,2-bal.rep`) that you can use for initial debugging.
- *Use the `mdriver -v` and `-V` options.* The `-v` option will give you a detailed summary for each trace file. The `-V` will also indicate when each trace file is read, which will help you isolate errors, as well as give a preview of your score.
- *Understand every line of the `malloc` implementation we provided before starting.* The code we have provided you is inspired by the textbook, which has a detailed discussion of a first-fit implicit list allocator. Don't start working on your allocator until you understand everything about the simple implicit list allocator.
- *Encapsulate your pointer arithmetic in `C` preprocessor macros.* Pointer arithmetic in memory managers is confusing and error-prone because of all the casting that is necessary. You can reduce the complexity significantly by writing macros for your pointer operations. See the textbook for examples.
- *Use a profiler.* You may find the `gprof` tool helpful for optimizing performance.
- *Use the `debug build` target in the `Makefile`:* By running `make debug` you can build the program without any optimizations, making it easier to follow in `gdb`. But be sure to build with just `make` before testing for performance.
- *Start early!* It is possible to write an efficient `malloc` package with a few pages of code. However, we can guarantee that it will be some of the most difficult and sophisticated code you have written so far in your career. So start early, and good luck!